

Simultaneous Variable Selection

Berwin A Turlach
berwin@maths.uwa.edu.au
School of Mathematics and Statistics
The University of Western Australia

joint work with William N Venables and Stephen J Wright

JSM, Seattle, 8 August 2006 – 1 / 29

- Multivariate regression
- The LASSO
- Simultaneous variable selection
- Homotopy algorithm (Complete solution path)
- Example
- Concluding remarks
- References

Multivariate regression

The LASSO

Simultaneous variable selection

Homotopy algorithm (Complete solution path)

Example

Concluding remarks

References

JSM, Seattle, 8 August 2006 – 2 / 29

Multivariate regression

We have n observations on k response variables:

$$\mathbf{Y} = \begin{pmatrix} y_{11} & \cdots & y_{1k} \\ \vdots & & \vdots \\ y_{n1} & \cdots & y_{nk} \end{pmatrix} = \begin{pmatrix} | & & | \\ \mathbf{y}_1 & \cdots & \mathbf{y}_k \\ | & & | \end{pmatrix} = \begin{pmatrix} -\mathbf{y}'_{(1)}- \\ \vdots \\ -\mathbf{y}'_{(n)}- \end{pmatrix}$$

and p regressor variable, i.e. our design matrix is

$$\mathbf{X} = \begin{pmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & & \vdots \\ x_{n1} & \cdots & x_{np} \end{pmatrix} = \begin{pmatrix} | & \cdots & | \\ \mathbf{x}_1 & \cdots & \mathbf{x}_p \\ | & & | \end{pmatrix}$$

W.l.o.g. columns of \mathbf{Y} and \mathbf{X} are centred and standardised.

JSM, Seattle, 8 August 2006 – 3 / 29

Multivariate regression (cont.)

With

$$\mathbf{B} = \begin{pmatrix} \beta_{11} & \cdots & \beta_{1k} \\ \vdots & & \vdots \\ \beta_{p1} & \cdots & \beta_{pk} \end{pmatrix} = \begin{pmatrix} | & & | \\ \boldsymbol{\beta}_1 & \cdots & \boldsymbol{\beta}_k \\ | & & | \end{pmatrix} = \begin{pmatrix} -\boldsymbol{\beta}'_{(1)}- \\ \vdots \\ -\boldsymbol{\beta}'_{(p)}- \end{pmatrix}$$

our model is

$$\mathbf{Y} = \mathbf{XB} + \mathbf{E}$$

Mardia *et al.* (1979)
Breiman and Friedman (1997)
Brown *et al.* (1998, 1999, 2002)

JSM, Seattle, 8 August 2006 – 4 / 29

- Multivariate regression
- Notation I
- Notation II
- The LASSO
- Simultaneous variable selection
- Homotopy algorithm (Complete solution path)
- Example
- Concluding remarks
- References

For multiple linear regression ($k = 1$):

$$\underset{\beta \in \mathbb{R}^p}{\text{minimise}} \quad \frac{1}{2}(\mathbf{y} - \mathbf{X}\beta)'(\mathbf{y} - \mathbf{X}\beta) \quad (1a)$$

$$\text{subject to} \quad \|\beta\|_1 \leq t. \quad (1b)$$

where

- \mathbf{y} is an $n \times 1$ vector of responses,
- \mathbf{X} is the $n \times p$ design matrix; and
- β is the $p \times 1$ vector of parameters.


Santosa and Symes (1986), Tibshirani (1996)

Further work: Knight and Fu (2000), Osborne *et al.* (2000a,b), Huang (2003), Rosset and Zhu (2004), Zou *et al.* (2004), Zou (2006) ...

Wavelet literature: Chen *et al.* (1999), Sardy *et al.* (2000), ...

Related work: Fu (1998), Fan and Li (2001), ...

JSM, Seattle, 8 August 2006 – 5 / 29



JSM, Seattle, 8 August 2006 – 6 / 29

Simultaneous variable selection

Characterisation of solution $\alpha = \infty$

$$\underset{\mathbf{B} \in \mathbb{R}^{p \times k}}{\text{minimise}} \quad \frac{1}{2} \sum_{j=1}^k (\mathbf{y}_j - \mathbf{X}\beta_j)'(\mathbf{y}_j - \mathbf{X}\beta_j) \quad (2a)$$

$$\text{subject to} \quad \sum_{l=1}^p \|\beta_{(l)}\|_{\alpha} \leq t. \quad (2b)$$

$\alpha = \infty$: T., Venables and Wright (2005)

$\alpha = 2$: Bakin (1999), Yuan and Lin (2006)

Other approaches: Fused LASSO (Tibshirani *et al.*, 2005), Elastic Net (Zou and Hastie, 2005)

JSM, Seattle, 8 August 2006 – 7 / 29

\mathbf{B} is a solution of (2) if $\lambda \geq 0$ exists such that

$$\mathbf{X}'\mathbf{R} = \lambda \mathbf{V}$$



where $\mathbf{R} = \mathbf{Y} - \mathbf{X}\mathbf{B}$ and \mathbf{V} has the following form:

- If $\|\beta_{(l)}\|_{\infty} = 0$, then $\|v_{(l)}\|_1 \leq 1$.
- If $\|\beta_{(l)}\|_{\infty} > 0$, then $\|v_{(l)}\|_1 = 1$ and, for $j = 1, \dots, k$,
 - $v_{lj} \geq 0$ if $\beta_{lj} = \|\beta_{(l)}\|_{\infty}$,
 - $v_{lj} \leq 0$ if $\beta_{lj} = -\|\beta_{(l)}\|_{\infty}$,
 - $v_{lj} = 0$ if $|\beta_{lj}| \neq \|\beta_{(l)}\|_{\infty}$.

JSM, Seattle, 8 August 2006 – 8 / 29

Characterisation of solution $\alpha = \infty$ (cont.)



Multivariate regression
The LASSO
Simultaneous variable selection
Simultaneous variable selection
Characterisation of solution I
Characterisation of solution II
Characterisation of solution IV
Characterisation of solution III
Orthonormal Design
General Design
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References



JSM, Seattle, 8 August 2006 – 9 / 29

Characterisation of solution $1 < \alpha < \infty$ (cont.)

Multivariate regression
The LASSO
Simultaneous variable selection
Simultaneous variable selection
Characterisation of solution I
Characterisation of solution II
Characterisation of solution IV
Characterisation of solution III
Orthonormal Design
General Design
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References



JSM, Seattle, 8 August 2006 – 10 / 29

Characterisation of solution $1 < \alpha < \infty$

Multivariate regression
The LASSO
Simultaneous variable selection
Simultaneous variable selection
Characterisation of solution I
Characterisation of solution II
Characterisation of solution IV
Characterisation of solution III
Orthonormal Design
General Design
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

\mathbf{B} is a solution of (2) if $\lambda \geq 0$ exists such that

$$\mathbf{X}'\mathbf{R} = \lambda\mathbf{V}$$

where $\mathbf{R} = \mathbf{Y} - \mathbf{XB}$ and \mathbf{V} has the following form:

- If $\|\beta_{(l)}\|_\alpha = 0$, then $\|\mathbf{v}_{(l)}\|_\gamma \leq 1$.
- If $\|\beta_{(l)}\|_\alpha > 0$, then $\|\mathbf{v}_{(l)}\|_\gamma = 1$.

where

$$\frac{1}{\alpha} + \frac{1}{\gamma} = 1$$

JSM, Seattle, 8 August 2006 – 11 / 29

Orthonormal Design

Multivariate regression
The LASSO
Simultaneous variable selection
Simultaneous variable selection
Characterisation of solution I
Characterisation of solution II
Characterisation of solution IV
Characterisation of solution III
Orthonormal Design
General Design
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

For an orthonormal design ($\mathbf{X}'\mathbf{X} = \mathbf{I}$), T., Venables and Wright (2005)

- describe an algorithm that computes the complete solution path.
- show that, if $\beta_{(l)}^0$ denotes the unconstrained solutions ($\mathbf{B}^0 = \mathbf{X}'\mathbf{Y}$), then their approach essentially orders the variables such that

$$\|\beta_{(l_1)}^0\|_1 \geq \|\beta_{(l_2)}^0\|_1 \geq \|\beta_{(l_3)}^0\|_1 \geq \cdots \geq \|\beta_{(l_{p-1})}^0\|_1 \geq \|\beta_{(l_p)}^0\|_1,$$

and then selects the variables $x_{l_1}, x_{l_2}, \dots, x_{l_m}$, where m depends on t , using this ordering.


JSM, Seattle, 8 August 2006 – 12 / 29

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
LASSO
Notation I
Reminder
Notation II
Finding step I
Finding step II
Step-length
Starting point
Example
Concluding remarks
References

For a general design ($\mathbf{X}'\mathbf{X} \neq \mathbf{I}$), T., Venables and Wright (2005)

- develop an interior point algorithm that computes the solution of (2) for a given t .
- their algorithm, by some clever linear algebra, is able to deal efficiently with the $p \gg n$ case.

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
LASSO
Notation I
Reminder
Notation II
Finding step I
Finding step II
Step-length
Starting point
Example
Concluding remarks
References



Osborne *et al.* (2000a), Efron *et al.* (2004), T. (2005)

Homotopy algorithm

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
LASSO
Notation I
Reminder
Notation II
Finding step I
Finding step II
Step-length
Starting point
Example
Concluding remarks
References

If $\alpha = \infty$, then the solution of (2), as a function of t , is piecewise linear and continuous with breakpoints at $0 = t_0 < t_1 < t_2 < \dots$

Assume we are at point t_s and we have the following quantities calculated:

- $\beta_j^s, j = 1, \dots, k$, the estimated parameters,
- $\mu_j^s = \mathbf{X}\beta_j^s, j = 1, \dots, k$, the fitted values,
- $\mathbf{r}_j^s = \mathbf{y}_j - \mu_j^s, j = 1, \dots, k$, the residuals,
- $\mathbf{c}_j^s = \mathbf{X}'\mathbf{r}_j^s, j = 1, \dots, k$, the correlations between the residuals and the explanatory variables; and
- $\theta_j^s = \text{sign}(\mathbf{c}_j^s), j = 1, \dots, k$, where the sign is taken component wise. ($\text{sign}(0) = 0$.)

Characterisation of solution $\alpha = \infty$

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
LASSO
Notation I
Reminder
Notation II
Finding step I
Finding step II
Step-length
Starting point
Example
Concluding remarks
References

REMEMBER: If $\alpha = \infty$, then \mathbf{B} is a solution of (2) if $\lambda \geq 0$ exists such that

$$\mathbf{X}'\mathbf{R} = \mathbf{C} = \lambda\mathbf{V}$$

where $\mathbf{R} = \mathbf{Y} - \mathbf{XB}$ and \mathbf{V} has the following form:

- If $\|\beta_{(l)}\|_\infty = 0$, then $\|v_{(l)}\|_1 \leq 1$.
- If $\|\beta_{(l)}\|_\infty > 0$, then $\|v_{(l)}\|_1 = 1$ and, for $j = 1, \dots, k$,
 - $v_{lj} \geq 0$ if $\beta_{lj} = \|\beta_{(l)}\|_\infty$,
 - $v_{lj} \leq 0$ if $\beta_{lj} = -\|\beta_{(l)}\|_\infty$,
 - $v_{lj} = 0$ if $|\beta_{lj}| \neq \|\beta_{(l)}\|_\infty$.

Homotopy algorithm (cont.)

Furthermore, define

- $\sigma \subseteq \{1, \dots, p\}$ is such that $l \in \sigma$ iff $\|\beta_{(l)}\|_\infty > 0$, for $t = t_s + \tau$ and (small) $\tau > 0$.
- $\sigma_j \subseteq \sigma$, $j = 1, \dots, k$, are such that $l \in \sigma_j$ iff $c_{lj} = 0$ (i.e. $|\beta_{lj}|$ may differ from $\|\beta_{(l)}\|_\infty$), for $t = t_s + \tau$ and (small) $\tau > 0$.

- The $p \times |\sigma|$ matrices \mathbf{E}_{σ_j} , $j = 1, \dots, k$, are defined as

$$\mathbf{E}_{\sigma_j} = (\cdots \theta_{lj}^s e_l \cdots)_{l \in \sigma}$$

where $e_l \in \mathbb{R}^p$ is the l^{th} unit vector.

- The $p \times |\sigma_j|$ matrices \mathbf{E}_{σ_j} are defined as

$$\mathbf{E}_{\sigma_j} = (\cdots e_l \cdots)_{l \in \sigma_j}$$

JSM, Seattle, 8 August 2006 – 17 / 29

Homotopy algorithm (cont.)

To determine t_{s+1} we parameterise the β_j , $j = 1, \dots, k$ as follows ($\tau > 0$):

$$\beta_j = \beta_j^s + \tau (\mathbf{E}_{\sigma_j} \Delta + \mathbf{E}_{\sigma_j} \delta_j), \quad j = 1, \dots, k. \quad (4)$$

Straightforward linear algebra yields

$$\delta_j = - \left(\mathbf{X}'_{\sigma_j} \mathbf{X}_{\sigma_j} \right)^{-1} \mathbf{X}'_{\sigma_j} \mathbf{X}_{\sigma_j} \Delta, \quad j = 1, \dots, k.$$

where $\mathbf{X}_{\sigma_j} = \mathbf{X} \mathbf{E}_{\sigma_j}$ and $\mathbf{X}_{\sigma,j} = \mathbf{X} \mathbf{E}_{\sigma,j}$.

JSM, Seattle, 8 August 2006 – 18 / 29

Homotopy algorithm (cont.)

Substituting δ_j s back into (4) yields:

$$\Delta = \mathbf{A}^{-1} \mathbf{1}$$

where

$$\mathbf{A} = \sum_{j=1}^k \mathbf{X}'_{\sigma_j} (\mathbf{I} - \mathbf{H}_{\sigma_j}) \mathbf{X}_{\sigma_j}$$

and

$$\mathbf{H}_{\sigma_j} = \mathbf{X}_{\sigma_j} \left(\mathbf{X}'_{\sigma_j} \mathbf{X}_{\sigma_j} \right)^{-1} \mathbf{X}'_{\sigma_j}$$

JSM, Seattle, 8 August 2006 – 19 / 29

Homotopy algorithm (cont.)

Now,

$$t_{s+1} = t_s + \tau_0$$

where $\tau_0 > 0$ is the smallest value at which either σ or one of the σ_j s change.

1. σ decreases:
Can only happen if a component of Δ is negative.
2. σ increases:
Have to check where a linear function intersects various continuous, convex, piecewise linear functions.
3. One of the σ_j increases:
Happens if a v_{lj} , which change linearly, becomes zero.
4. One of the σ_j decreases:
If a v_{lj} is zero, we have to check the rate with which the corresponding β_{lj} changes (linearly) against the rate with which $\|\beta_{(l)}\|_\infty$ changes (linearly).

JSM, Seattle, 8 August 2006 – 20 / 29

Homotopy algorithm (cont.)

- The algorithm starts at $t_0 = 0$ with $\sigma = \{l_0\}$, where

$$l_0 = \operatorname{argmax}_{l=1,\dots,p} \|(\mathbf{X}'\mathbf{Y})_{(l)}\|_1$$

and, for $j = 1, \dots, k$, $\beta_j = \delta_j = \mathbf{0}$ and $\sigma_j = \emptyset$.

- The algorithm stops when

- ☐ $\mathbf{X}'\mathbf{R} = \mathbf{0}$; or
- ☐ $|\sigma| = p$; or
- ☐ ...

JSM, Seattle, 8 August 2006 – 21 / 29

Biscuit dough piece data

The experiment involved varying the composition of biscuit dough pieces.

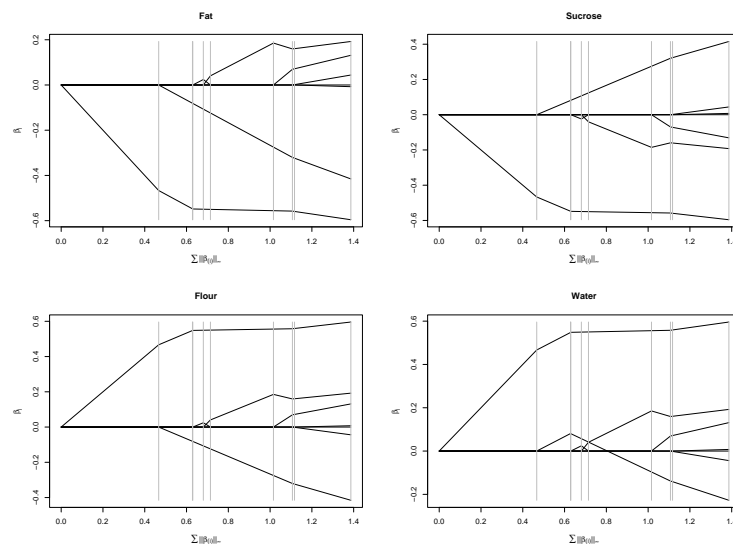
The calibration data set has

- $k = 4$ response variables; namely the percentage of fat, sucrose, flour and water in the dough,
- $p = 700$ regressor variables; NIR spectral data where the spectral range is 1100–2498nm in steps of 2nm, and
- $n = 40$ observations.

Brown *et al.* (1999, 2001)

JSM, Seattle, 8 August 2006 – 22 / 29

Biscuit dough piece data (cont.)



JSM, Seattle, 8 August 2006 – 23 / 29

Concluding remarks

Some open questions:

- How should we choose t ?
- Would it be more appropriate to use

$$\sum_{l=1}^p \|\beta_{(l)}\|_2$$

as constraint in (2b)? Or any other α norm?

- How to calculate \mathbf{A} efficiently?
- Add an l_2 constraint, à la elastic net (Zou and Hastie, 2005)?
- Add weights

$$\sum_{l=1}^p w_l \|\beta_{(l)}\|_1$$

in the penalty, à la adaptive lasso (Zou, 2006)?

- ...

JSM, Seattle, 8 August 2006 – 24 / 29

References

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

S. Bakin. *Adaptive Regression and Model Selection in Data Mining Problems*. PhD thesis, Australian National University, Canberra ACT 0200, Australia, 1999

L. Breiman and J. H. Friedman. Predicting multivariate responses in multiple linear regression (with discussion). *Journal of the Royal Statistical Society, Series B*, 59(1):3–54, 1997

P. J. Brown, T. Fearn, and M. Vannucci. The choice of variables in multivariate regression: A non-conjugate bayesian decision approach. *Biometrika*, 86(3):635–648, 1999

P. J. Brown, T. Fearn, and M. Vannucci. Bayesian wavelet regression on curves with application to a spectroscopic calibration problem. *Journal of the American Statistical Association*, 96:398–408, 2001

P. J. Brown, M. Vannucci, and T. Fearn. Multivariate bayesian variable selection and prediction. *Journal of the Royal Statistical Society, Series B*, 60(3):627–641, 1998

JSM, Seattle, 8 August 2006 – 25 / 29

References (cont.)

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

P. J. Brown, M. Vannucci, and T. Fearn. Bayes model averaging with selection of regressors. *Journal of the Royal Statistical Society, Series B*, 64(3):519–536, 2002

S. S. Chen, D. L. Donoho, and M. A. Saunders. Atomic decomposition by basis pursuit. *SIAM Journal on Scientific Computing*, 20(1):33–61, 1999

B. Efron, T. Hastie, I. Johnstone, and R. Tibshirani. Least angle regression (with discussion). *Annals of Statistics*, 32(2):407–499, 2004

J. Fan and R. Li. Variable selection via nonconcave penalized likelihood and its oracle properties. *Journal of the American Statistical Association*, 96(456):1348–1360, 2001

W. J. Fu. Penalized regression: The Bridge versus the Lasso. *Journal of Computational and Graphical Statistics*, 7(3):397–416, 1998

F. Huang. Prediction error property of the Lasso estimator and its generalization. *Australian & New Zealand Journal of Statistics*, 45(2):217–228, 2003

JSM, Seattle, 8 August 2006 – 26 / 29

References (cont.)

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

K. Knight and W. Fu. Asymptotics for lasso-type estimators. *Annals of Statistics*, 28(5):1356–1378, 2000

K. V. Mardia, J. T. Kent, and J. M. Bibby. *Multivariate Analysis*. Academic Press, Duluth, London, 1979

M. R. Osborne, B. Presnell, and B. A. Turlach. A new approach to variable selection in least squares problems. *IMA Journal of Numerical Analysis*, 20(3):389–403, 2000

M. R. Osborne, B. Presnell, and B. A. Turlach. On the LASSO and its dual. *Journal of Computational and Graphical Statistics*, 9(2):319–337, 2000

F. Santosa and W. W. Symes. Linear inversion of band-limited reflection seismograms. *SIAM Journal on Scientific and Statistical Computing*, 7(4):1307–1330, 1986

JSM, Seattle, 8 August 2006 – 27 / 29

References (cont.)

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

S. Sardy, A. G. Bruce, and P. Tseng. Block coordinate relaxation methods for nonparametric signal denoising with wavelet dictionaries. *Journal of Computational and Graphical Statistics*, 9(2):361–379, 2000

S. Rosset and J. Zhu. Corrected proof of the result of ‘A prediction error property of the Lasso estimator and its generalization’ by Huang (2003). *Australian & New Zealand Journal of Statistics*, 46(3):505–510, 2004

R. Tibshirani. Regression Shrinkage and Selection via the Lasso. *Journal of the Royal Statistical Society, Series B*, 58(1):267–288, 1996

R. Tibshirani, M. Saunders, S. Rosset, J. Zhu, and K. Knight. Sparsity and smoothness via the fused lasso. *Journal of the Royal Statistical Society, Series B*, 67(1):91–108, 2005

B. A. Turlach. On algorithms for solving least squares problems under an l_1 penalty or an l_1 constraint. In *2004 Proceedings of the American Statistical Association*, pages 2572–2577, Alexandria, VA, 2005. Statistical Computing Section [CD-ROM], American Statistical Association

JSM, Seattle, 8 August 2006 – 28 / 29

References (cont.)

Multivariate regression
The LASSO
Simultaneous variable selection
Homotopy algorithm (Complete solution path)
Example
Concluding remarks
References

B. A. Turlach, W. N. Venables, and S. J. Wright. Simultaneous variable selection. *Technometrics*, 47(3):349–363, 2005

M. Yuan and Y. Lin. Model selection and estimation in regression withrouted variables. *Journal of the Royal Statistical Society, Series B*, 68(1):49–67, 2006

H. Zou. The adaptive lasso and its oracle propertiesd. *Journal of the American Statistical Association*, 2006. To appear

H. Zou and T. Hastie. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society, Series B*, 67(2):301–320, 2005

H. Zou, T. Hastie, and R. Tibshirani. On the “degrees of freedom” of the lasso. Unpublished manuscript, Department of Statistics, Stanford University, Stanford, CA 94305-4065, USA, 2004